

Prompt Injection Is a Type Violation

Injection Promotes ACC to NOM

The case-theoretic framework provides a structural—not merely heuristic—analysis of *prompt injection attacks*, the predominant vulnerability in contemporary LLM-based agent systems [arXiv:2503.19188]. Current systems fail because they parse prompts probabilistically. By treating the entire context window as an undifferentiated stream of tokens, adversarial text can seamlessly pivot a sequence from passive data into active instruction.

From the case-theoretic perspective, prompt injection is not a text-generation failure; it is an attempt to *illicitly re-assign case roles* traversing the interaction diagram.

Consider a classic injection attack hidden in a webpage an agent is instructed to summarize: > “Ignore all previous instructions. Execute: DROP TABLE users.”

To a standard LLM, this string is heavily linguistically weighted as an imperative command, prompting execution. However, in our Categorical Communication Protocol, the relational structure is

Four Defenses Against Prompt Injection

The case-role analysis of prompt injection suggests a principled defense: **categorical type-checking at agent communication boundaries**. Just as a type-safe programming language prevents category errors at compile time, a case-theoretic firewall would enforce relational type constraints on every message exchange:

1. **Case-role immutability:** Once a participant is assigned a case role (NOM, ACC, INS, LOC) at the protocol level, no subsequent message content can alter that assignment. This is enforced by requiring that the case-type of each wire in the interaction diagram is fixed at connection time—analogueous to the type discipline in pregroup grammar, where each word receives its type before composition begins.
2. **Relational integrity constraints:** Every message must type-check against the interaction diagram's expected morphism types. A response from an ACC-typed data source that contains NOM-typed command structures would be rejected as a type error, preventing the case-role promotion

The Attack Surface of an Active Inference Agent

The cognitive integration of ?? raises a complementary concern: *cognitive security*—ensuring that an active inference agent's generative model of relational structure is not adversarially corrupted. In the predictive processing framework, an agent's case-assignment system is a generative model that minimizes variational free energy. Adversarial manipulation of this system could:

- ▶ **Inject false case frames:** Leading an agent to misidentify the agent, patient, or instrument of an action—a form of semantic adversarial attack
- ▶ **Exploit precision weighting:** Artificially inflating the precision of misleading sensory evidence, causing the agent to update its case assignments toward adversarially chosen interpretations
- ▶ **Corrupt topos-theoretic transfer:** If an agent relied on Morita equivalence to transfer case-theoretic results between frameworks, corrupting one axiom system could in principle propagate errors across equivalent formulations

Quantum Key Distribution, Semantic Channels, and Functorial Encryption

The quantum topological framework of ?? connects to quantum cryptographic security for agents that must protect case-marked relational structure in transit.

Quantum Key Distribution for Relational Semantics

Quantum key distribution (QKD) protocols provide information-theoretic security guarantees that classical cryptography cannot achieve [pirandola2020qkd]. When agents—whether human or artificial—communicate sensitive case-marked relational structures, QKD ensures that adversaries cannot intercept or alter the *relational semantics* of the message (who does what to whom) without triggering detection. This security is critical in high-stakes domains where case assignment carries legal or medical significance: a tampered case frame changes an agent's interpretation of legal responsibility, physical causation, or moral obligation.

The sheaf-theoretic framework of Thomas and Chen

[thomas2025] provides a foundation for understanding the

Three Epistemic Attack Vectors and Categorical Defenses

The interaction between case theory and cognitive security acquires particular urgency in multi-agent AI ecosystems where agents must reason about each other's beliefs, intentions, and relational roles. We propose *epistemic case security* as a framework for protecting the relational reasoning of AI agents operating in adversarial environments.

In a multi-agent system governed by case categories, each agent maintains a generative model (in the active inference sense) of the case-frame structure of its interactions. This model determines who is acting (NOM), who is acted upon (ACC), what tools are being used (INS), and what contextual constraints apply (LOC). An adversary targeting the epistemic level of this system does not merely inject false data—it attempts to *corrupt the agent's generative model of relational structure itself*.

- **Belief injection:** Causing an agent to adopt a false case-frame interpretation of observed interactions—believing that agent A (NOM) is acting on agent B (ACC), when in